

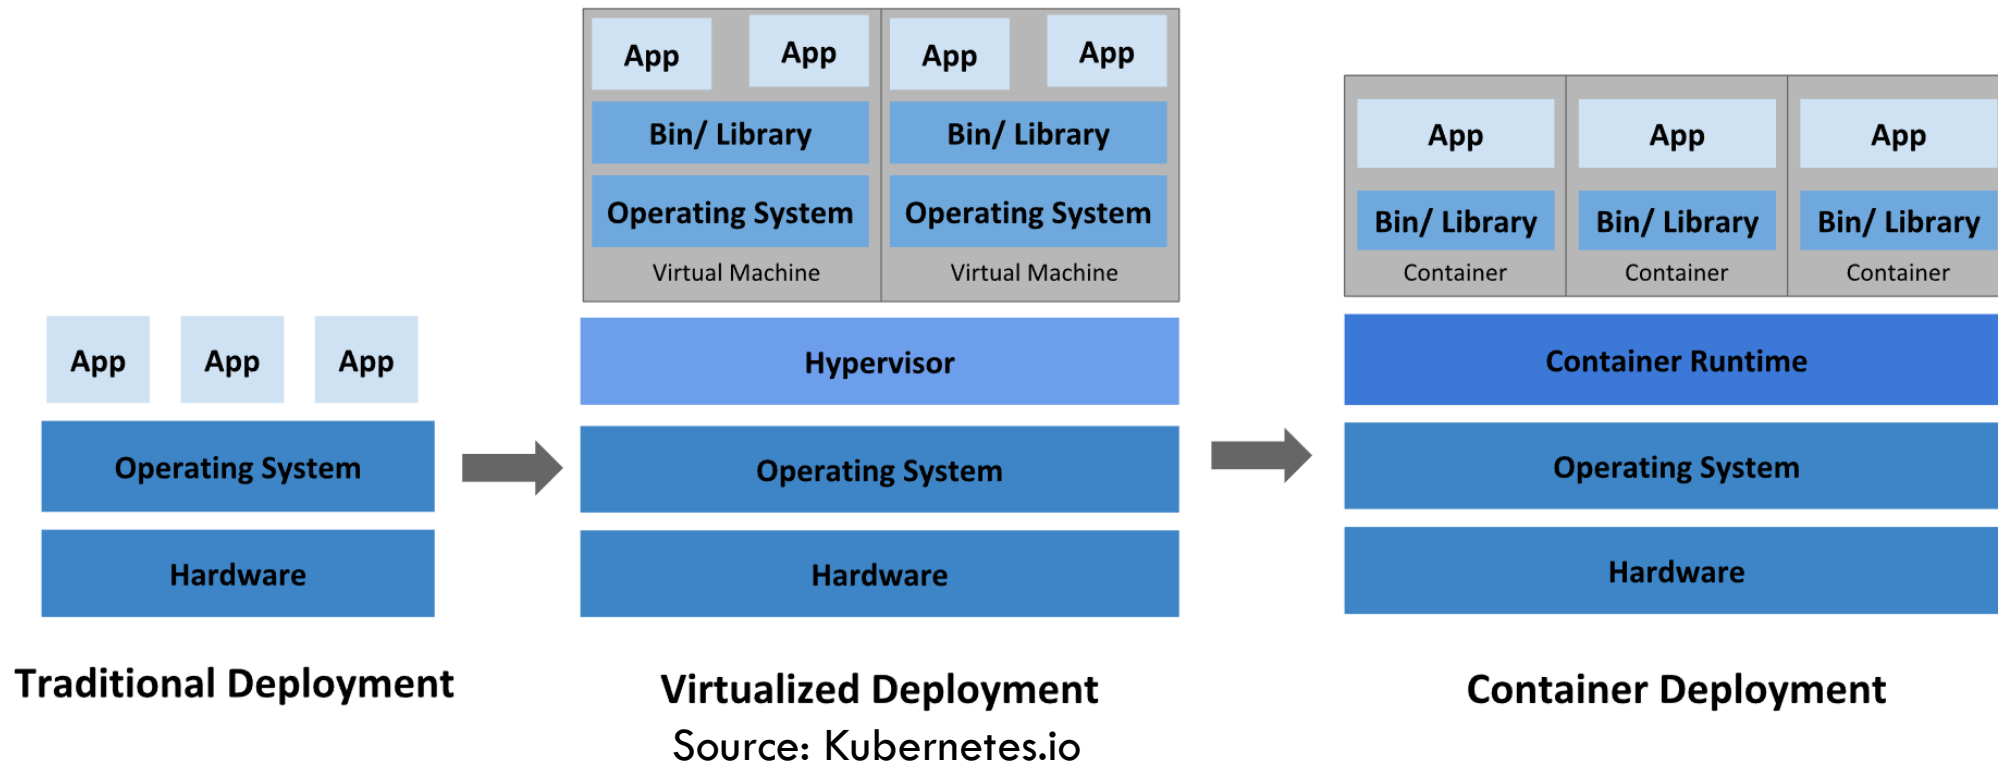
Lab 3 Part 2

A crash course on  **kubernetes**

Introduction

What is Kubernetes?


“Kubernetes, also known as K8s, is an open-source system for automating deployment, scaling, and management of containerized applications.”



Bare metal vs VM vs Container

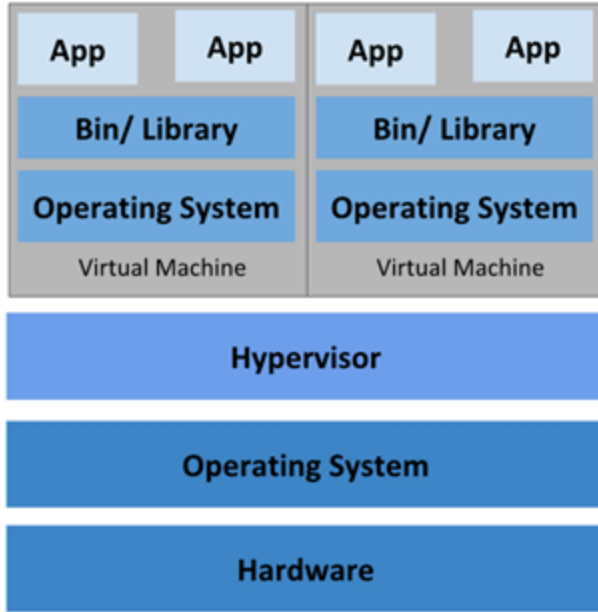
A traditional deployment runs applications on top of the OS, also known as running the applications on bare metal.

- ❖ Great efficiency and usage of resources
- ❖ Poor portability
- ❖ Resource sharing with no policies
 - ❖ Now, systemd allows you to put constraints on services
- ❖ Library sharing and incompatibilities
- ❖ Security issues



Traditional Deployment

Bare metal vs VM vs Container




Virtualized Deployment

A virtualized deployment runs group of applications on several VMs. All VMs are managed by a program (hypervisor)

- ❖ Too much overhead and performance loss
- ❖ Highly portable
- ❖ Hypervisors configure resource restrictions
- ❖ No library sharing. Each VM gets its one environment.
- ❖ Less security concerns

Bare metal vs VM vs Container



Container Deployment



A containerized deployment runs each application on a container. All containers are managed by a program (container runtime)

- ❖ Less overhead
- ❖ Highly portable
- ❖ Container runtime configure resource restrictions
- ❖ No library sharing. Each container gets its own environment.
- ❖ Less security concerns but...

Careful: One container only runs one process, no services/systemd inside.

What is Kubernetes cluster?

Let's explore the components of a k8s cluster




Source: Kubernetes.io


What is the control plane?

A k8s cluster is composed of two main components:

- ❖ Control plane:
 - ❖ Global services of the cluster.
 - ❖ Components:
 - ❖ etcd: HA key-value store
 - ❖ kube-scheduler: scheduling pods to nodes
 - ❖ kube-controller-manager: runs several controller processes such as node controller (check node status) or service controller (endpoint management)
 - ❖ kube-apiserver: frontend API to the users



What is a node?




A k8s cluster is composed of two main components:

- ❖ **Node:**

- ❖ Workers. They run Pods.
- ❖ They might be physical hosts or VMs.
- ❖ **Some essential components:**
 - ❖ kubelet: manage containers in Pods
 - ❖ kube-proxy: endpoints and forwarding for services
 - ❖ Container runtime: software that implements the Container Runtime Interface (CRI) to start/stop/monitor containers

But there is more...




A k8s cluster also needs:

- ❖ Control (and data) networks.
 - ❖ We want nodes to connect to the control plane to be managed.
 - ❖ We may need exchange of data among nodes.
- ❖ Addons: DNS (automatic service name resolution), dashboard (cool web UI), resource monitoring ...

How do we run apps?

What is a pod?

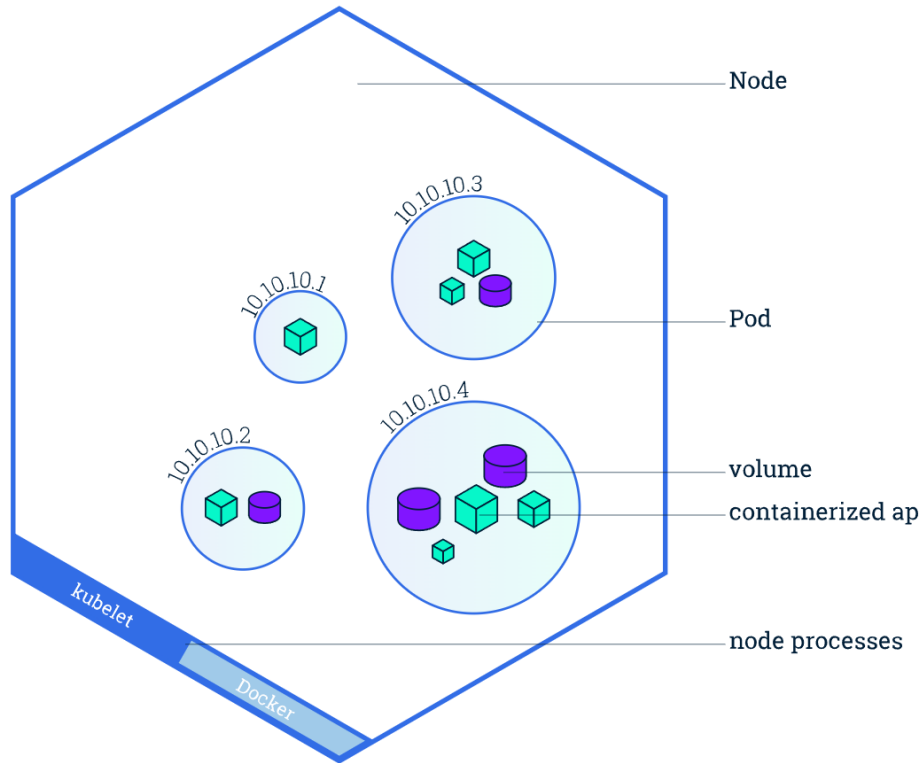


Pods are the smallest deployable units

- ❖ They are groups of containers
 - ❖ They shared storage and network resources
 - ❖ Containers should be tightly coupled
- ❖ Two models:
 - ❖ One pod per container: most common and simple to understand (pods = container)
 - ❖ One pod many containers: containers need to be co-located in the same node and need to share resources

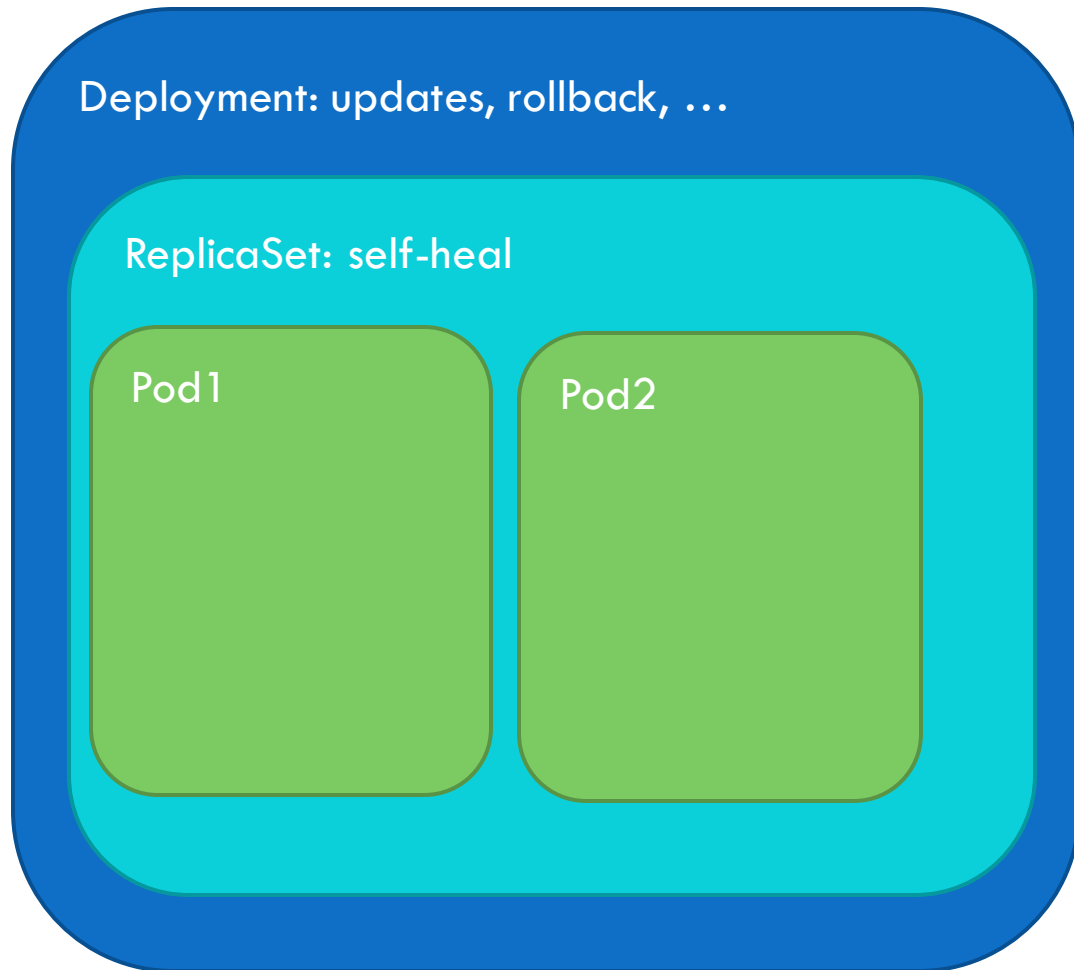
Creating a pod at least once in a lifetime

`kubectl apply -f pod.yaml`



```
apiVersion: v1
kind: Pod
metadata:
  name: examplepod
spec:
  volumes:
    - name: www_dir
      emptyDir: {}
  containers:
    - name: webserver
      image: nginx
      volumeMounts:
        - name: www_dir
          mountPath: /usr/share/nginx/html
    - name: updater
      image: debian
      volumeMounts:
        - name: www_dir
          mountPath: /html
      command: ["/bin/sh", "-c"]
      args:
        - while true; do
            date > /html/index.html;
            sleep 1;
          done
```

But nobody runs pods manually



In order to run Pods, we use high level abstractions:

- ❖ ReplicaSet: specifies that a precise number of Pods is running. Low-level compared to deployments.
- ❖ Deployment: similar but declarative. We describe the state we want, and the control plane creates/destroys ReplicaSets and Pods. **Recommended over ReplicaSets to create services**

How do we access apps?

What is a service?

By default, each Pod has a unique IP address in the internal network of k8s.

- ❖ However, since Pods are ephemeral, these IP addresses are dynamical
- ❖ In order to access a Pod, typically you should use a name.
- ❖ As an abstraction, services are used to refer to a Pod. This reference will be automatically handle by the cluster, even though Pods may die and respawn.

What is a service?

Services are abstraction to access applications, kind of providing an static way of accessing and exposing applications to the exterior or other Pods.

❖ There are three types:

- ❖ ClusterIP: internal IP address for the service. Only works for other Pods in the cluster.
- ❖ NodePort: nodes expose to the exterior the service by doing port forwarding
- ❖ Load Balancer: service is deployed through an existing load balancer. Usually, requires DNS services and it usually works only for L7 (http/https)

What is a ingress controller?

Ingress controllers allow to expose application to the exterior with state-of-the-art standards such as custom domain name or encryption (TLS).

- ❖ Typically, they just receive request and forward them to existing services of the cluster.
- ❖ Work on top well-known industrial HTTP/TCP/UDP servers
 - ❖ Nginx (HTTP, HTTPS)
 - ❖ Haproxy (TCP, UDP)
- ❖ For our case, we are not going to use them since there are more oriented towards web applications.

The kubectl CLI

How to manage objects in k8s?

List them

`kubectl get <type of object>`

```
kubectl get services
kubectl get pods
kubectl get replicaset
kubectl get nodes
Kubectl get deployments
```

You may run everything in isolated
namespaces

```
kubectl get services --all-namespaces
kubectl get pods -n spark
```

Extended information

`kubectl describe <type of object> <name>`

```
kubectl describe services
kubectl describe pods/nginx
kubectl describe -n spark replicaset
kubectl get nodes node1
```

Logs

`kubectl logs <podname>`

```
kubectl logs pods/nginx
```

Assignment

Exercise 1

Launch minikube on your computer.

- ❖ For the laboratories, follow the instructions on moodle
- ❖ For the rest, follow the guide at <https://minikube.sigs.k8s.io/docs/start/>

Follow a tutorial to check that it works.

- ❖ Try modifying some of the provided YAML files to try some options
- ❖ Link for the tutorial <https://kubernetes.io/docs/tutorials/stateless-application/guestbook/>

Exercise 2

Write a base Dockerfile based on ubuntu 22.04

- ❖ Install OpenJDK 11
- ❖ Install Hadoop (download and extract in /opt)
- ❖ Install Spark (download and extract in /opt)

Use that base Dockerfile to build the worker and master worker

Exercise 3

Write two deployments

1. For the master
2. For two workers

Write a service for the master exposing port 7077 for both internal usage and external access

Deploy both in the cluster and try to run a PySpark example on the host

- ❖ It should connect to localhost:7077
- ❖ Driver host should be the external IP address of your computer
 - ❖ Why? We need the Spark cluster to be able to find the PySpark program. Otherwise, the Driver (and the PySpark program) should run as a pod inside k8s.

Exercise 4

Scale up the number of workers to 3.

Check that the Spark dashboard shows all nodes.

Destroy the Spark cluster.

Can two Spark clusters co-exist on the same k8s infrastructure?

(Optional / Project ideas) What about storage for Spark on k8s? Ceph?

Material to submit

- ❖ You must write a report answering the questions proposed in each exercise, plus the requested files. Submit a zip file through Moodle. Check submission date in Moodle (deadline is until 11:59 pm of that date).
- ❖ From exercise 1:
 - ❖ Nothing
- ❖ From exercise 2:
 - ❖ Dockerfiles for creating the images
 - ❖ Explain how to build the images
- ❖ From exercise 3:
 - ❖ YAML file/s
 - ❖ Commented version of the YAML file
 - ❖ Explain how to launch the Deployments and Services on the cluster.
- ❖ From exercise 4:
 - ❖ Explain how to scale up and down the number of workers
 - ❖ Using kubectl apply and the YAML file
 - ❖ Using just kubectl
 - ❖ Explain how to destroy the Spark cluster
 - ❖ Explain how two Spark clusters can be deployed on the same infrastructure
 - ❖ (Extra credit/Project ideas) What about storage?